

OpenCL State of the Nation

Neil Trevett | Khronos President
NVIDIA Vice President Developer Ecosystem
OpenCL Working Group Chair
ntrevett@nvidia.com | [@neilt3d](https://twitter.com/neilt3d)
Toronto, May 2017



Topics

The Good

The amazing
progress of OpenCL

1

The Bad

Lessons Learned from
the first eight years

2

The Exciting

Where do we go
from here?

3

OpenCL 2.2 Finalized Here at IWOCL!



OpenCL 1.2

Becomes
industry
baseline

OpenCL 2.0
Enables new
class of
hardware

SVM
Generic Addresses
On-device dispatch

OpenCL 2.1
SPIR-V 1.0

SPIR-V in Core

Kernel Language Flexibility

OpenCL 2.2
SPIR-V 1.2

OpenCL C++ Kernel Language

Static subset of C++14

Templates and Lambdas

SPIR-V 1.2
OpenCL C++ support

Pipes

Efficient device-scope communication between kernels

Code Generation Optimizations

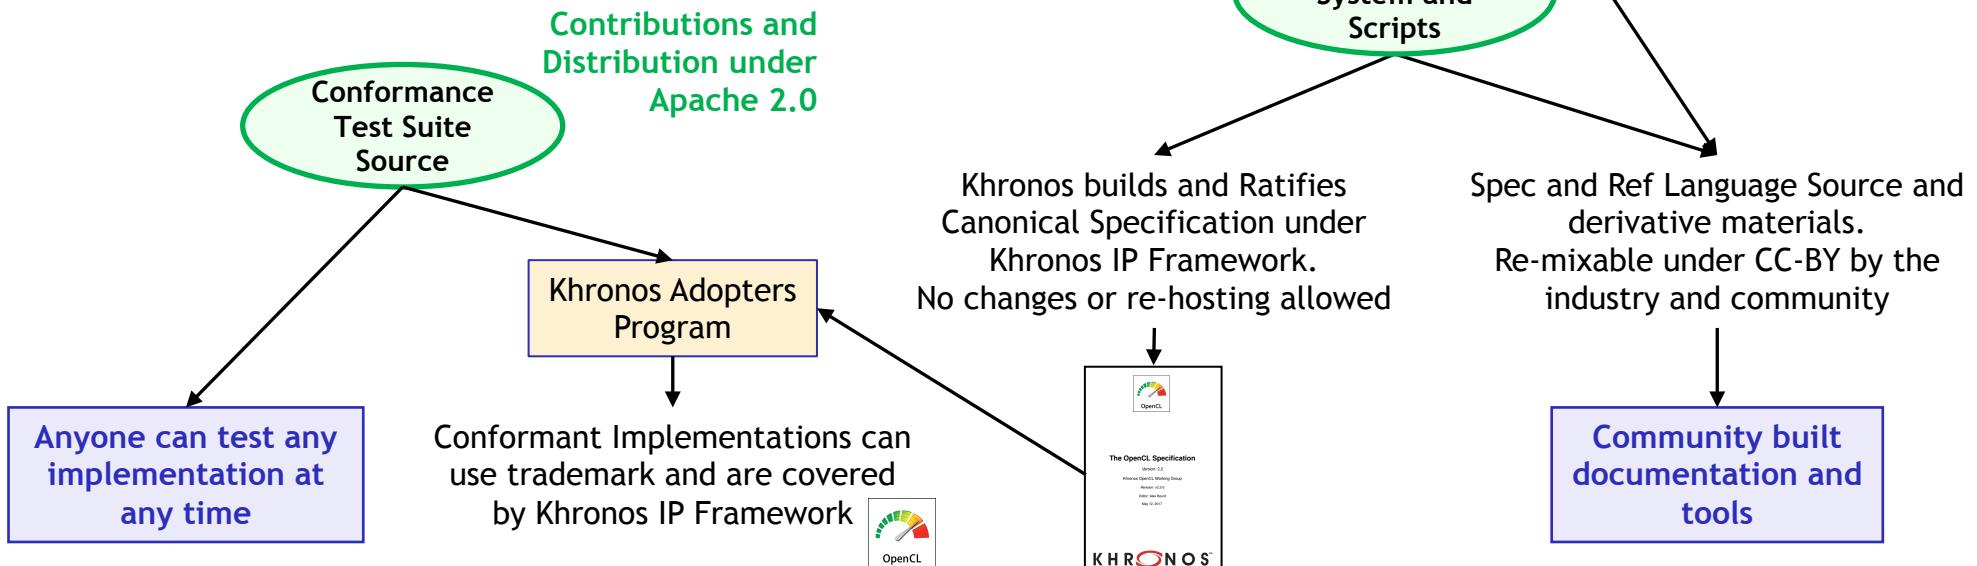
- Specialization constants at SPIR-V compilation time
 - Constructors and destructors of program scope global objects
 - User callbacks can be set at program release time

<https://www.khronos.org/opencl/>



New Open Source Engagement Model

- Khronos is open sourcing specification sources, conformance tests, tools
 - Merge requests welcome from the community (subject to review by OpenCL working group)
- Deeper Community Enablement
 - Mix your own documentation!
 - Contribute and fix conformance tests
 - Fix the specification, headers, ICD etc.
 - Contribute new features (carefully)



Shout Out to University of Windsor

The [Windsor Testing Framework](#), also released today, enables developers to quickly install and configure the OpenCL Conformance Test Suite on their own systems.



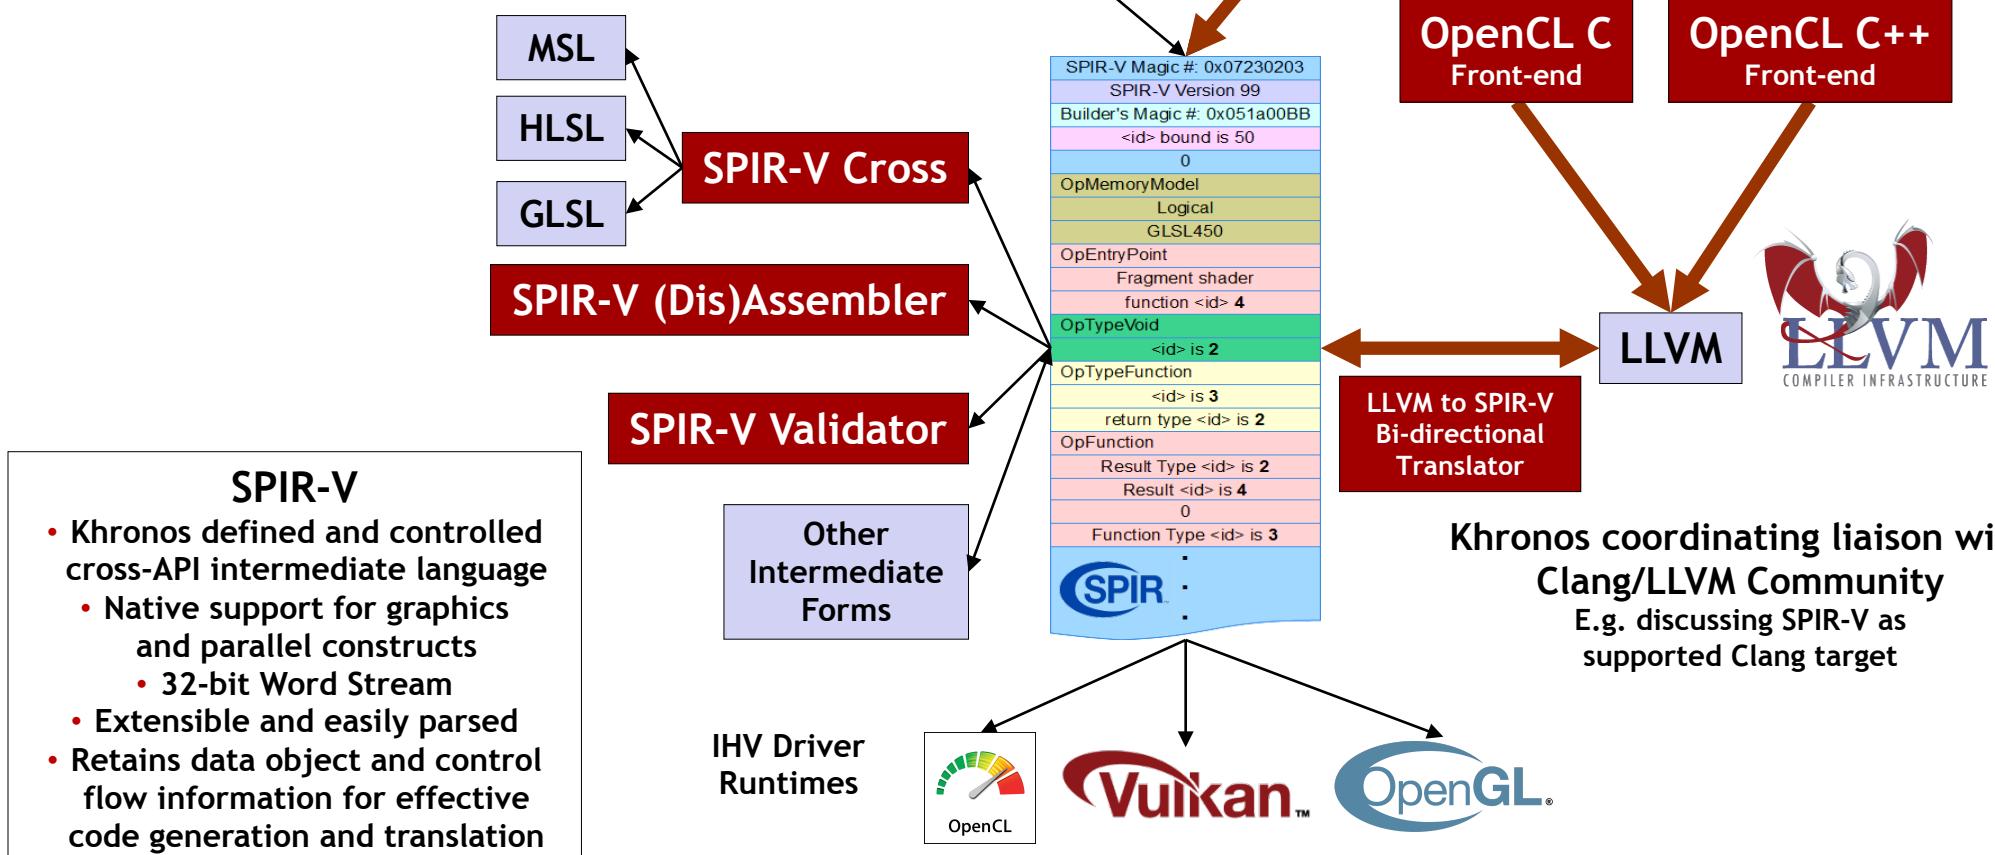
University
of Windsor

SPIR-V Ecosystem



Khronos has open sourced these tools and translators

<https://github.com/KhronosGroup/SPIRV-Tools>



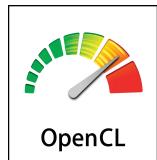
SYCL Ecosystem

- Single-source heterogeneous programming using STANDARD C++
 - Use C++ templates and lambda functions for host & device code
 - Layered over OpenCL
- Fast and powerful path for bring C++ apps and libraries to OpenCL
 - C++ Kernel Fusion - better performance on complex software than hand-coding
 - Halide, Eigen, Boost.Compute, SYCLBLAS, SYCL Eigen, SYCL TensorFlow, SYCL GTX
 - triSYCL, ComputeCpp, VisionCpp, ComputeCpp SDK ...
- More information at <http://sycl.tech>

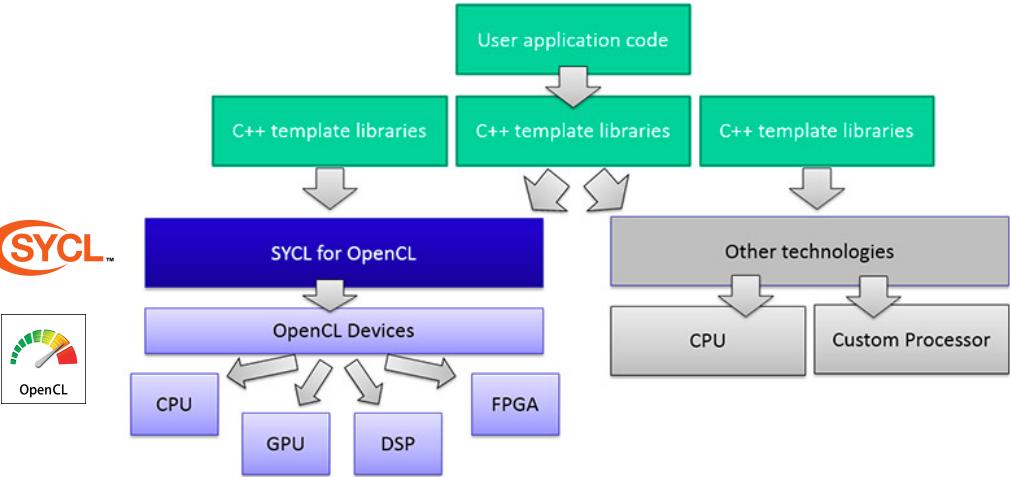
Developer Choice

The development of the two specifications are aligned so code can be easily shared between the two approaches

C++ Kernel Language
Low Level Control
'GPGPU'-style separation of device-side kernel source code and host code



Single-source C++
Programmer Familiarity
Approach also taken by C++ AMP and OpenMP



OpenCL Adoption

- 100s of applications using OpenCL acceleration
 - Rendering, visualization, video editing, simulation, image processing
- Over 4,000 GitHub repositories using OpenCL
 - Tools, applications, libraries, languages
 - Up from 2,000 two years ago
- Multiple silicon and open source implementations
 - Increasingly used for embedded vision and neural network inferencing
- Khronos Resource Hub

<https://www.khronos.org/opencl/resources/opencl-applications-using-opencl>



Adobe



WOLFRAM



otoy

SONY



ppcl



GIMP
GNU Image Manipulation Program

Blackmagicdesign



CHAOSGROUP

A screenshot of a GitHub search results page for the query 'opencl'. The page shows 4,310 repository results. The top repository listed is 'arrayfire/arrayfire', described as a 'general purpose GPU library' in C++. The second repository listed is 'LWJGL/lwjgl3', described as a Java library for cross-platform access to native APIs. The GitHub interface includes a header with the search term 'opencl', navigation links for 'Repositories 4K', 'Code 1M', 'Commits 204K', 'Issues 32K', 'Wikis 4K', and 'Users 123', and a dropdown menu for sorting results by 'Most stars'.

OpenCL as Language/Library Backend

Caffe

Halide

C++ AMP
Accelerated Massive Parallelism
with Microsoft Visual C++™

SYCL™

aparapi

OpenCV

OpenACC®
DIRECTIVES FOR ACCELERATORS

TensorFlow

C++ based
Neural
network
framework

Language for
image
processing and
computational
photography

MulticoreWare
open source
project on
Bitbucket

Single
Source C++
Programming
for OpenCL

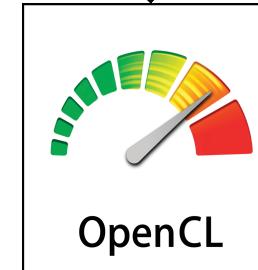
Java language
extensions
for
parallelism

Vision
processing
open source
project

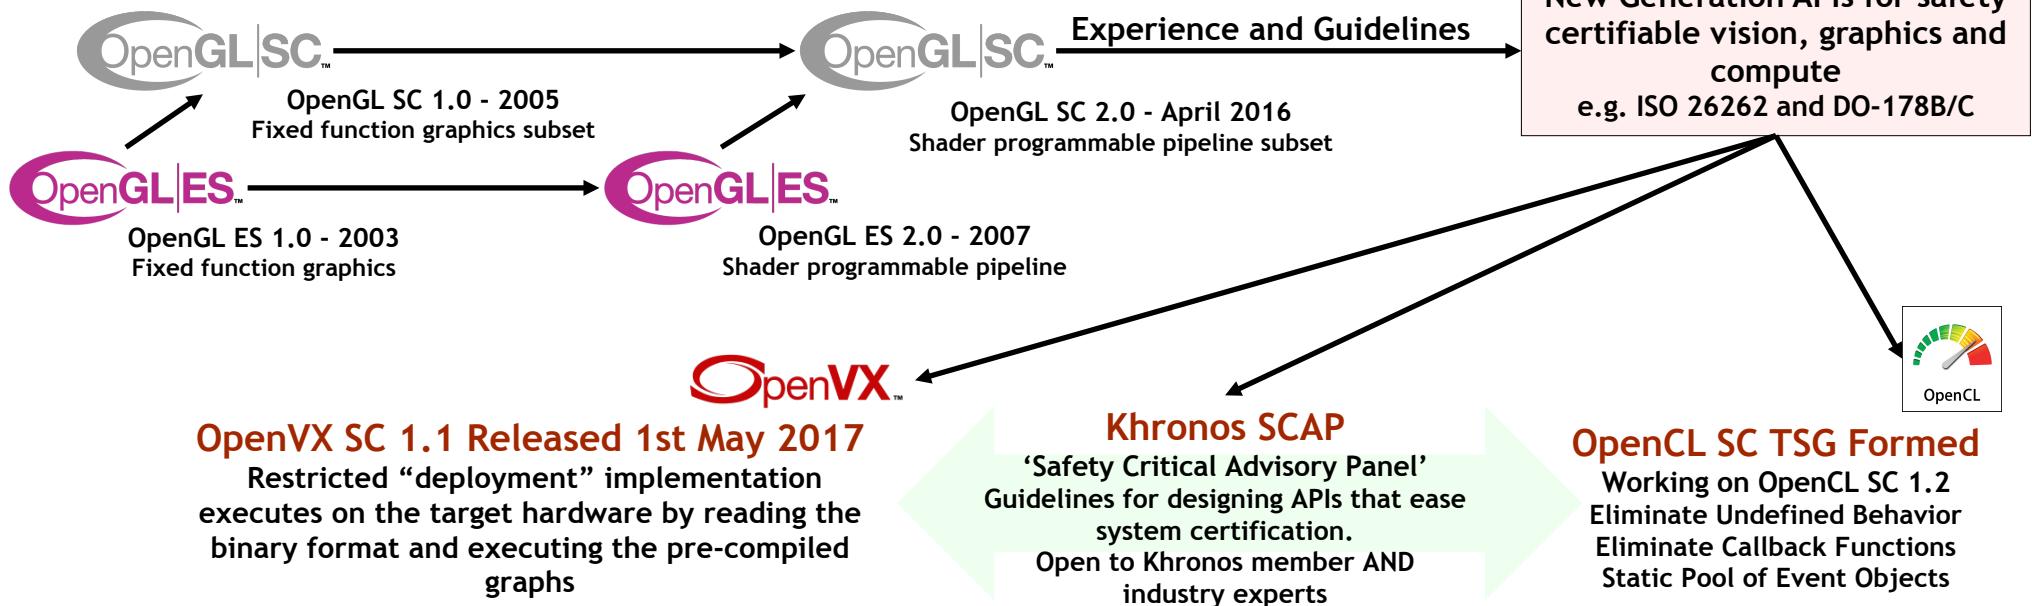
Compiler
directives for
Fortran,
C and C++

Open source
software library
for machine
learning

Hundreds of languages, frameworks
and projects using OpenCL to access
vendor-optimized, heterogeneous
compute runtimes



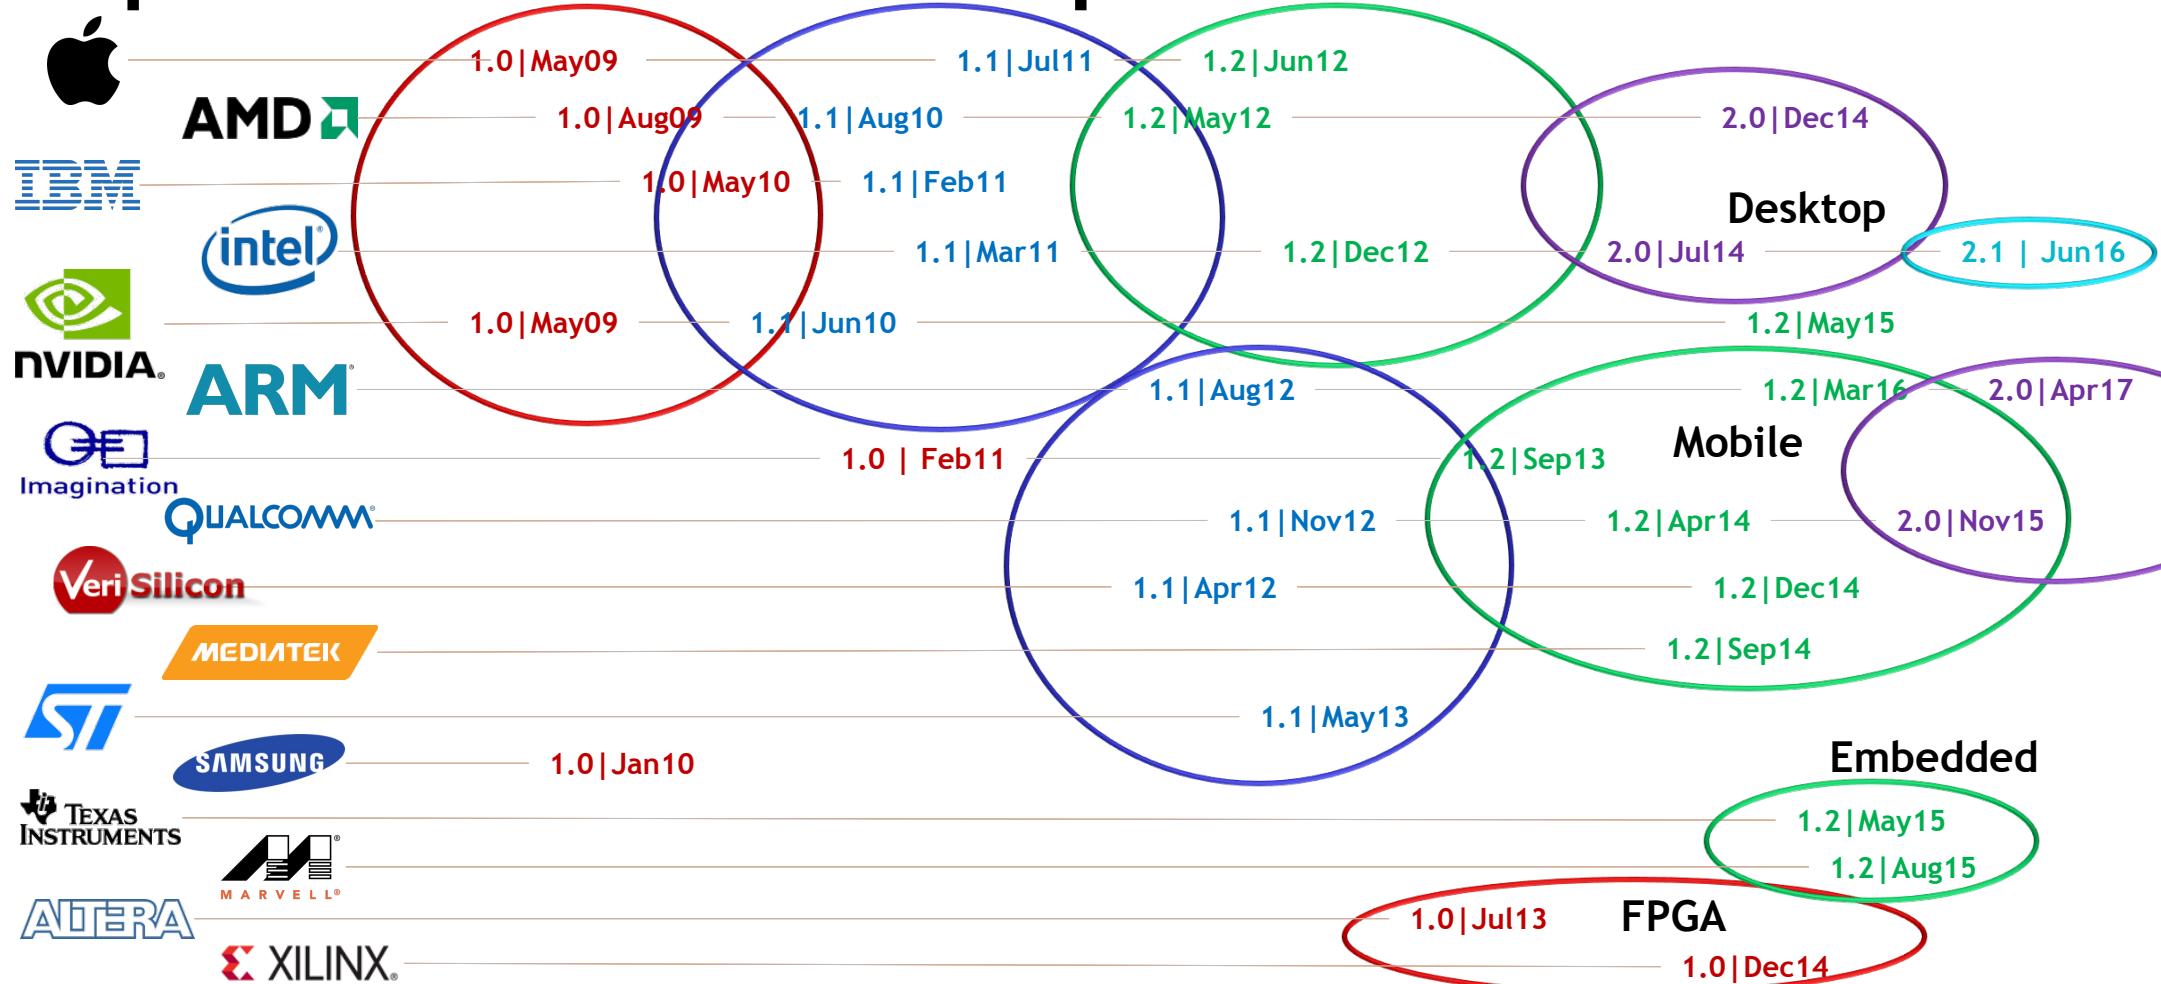
Safety Critical APIs



OpenVX SC 1.1 Released 1st May 2017

Restricted “deployment” implementation executes on the target hardware by reading the binary format and executing the pre-compiled graphs

OpenCL Conformant Implementations



Vendor timelines are
first conformant
submission for each spec
generation

Dec08
OpenCL 1.0
Specification

Jun10
OpenCL 1.1
Specification

Nov11
OpenCL 1.2
Specification

Nov13
OpenCL 2.0
Specification

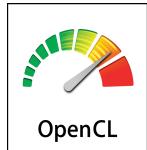
Nov15
OpenCL 2.1
Specification

OpenCL - 1000s Man Years Effort



Single Source C++ Programming

Full support for features in C++14-based Kernel Language



API and Language Specs

Brings C++14-based Kernel Language into core specification



Portable Kernel Intermediate Language

Support for C++14-based kernel language e.g.
constructors/destructors

3-component vectors
Additional image formats
Multiple hosts and devices
Buffer region operations
Enhanced event-driven execution
Additional OpenCL C built-ins
Improved OpenGL data/event interop

Device partitioning
Separate compilation and linking
Enhanced image support
Built-in kernels / custom devices
Enhanced DX and OpenGL Interop

Shared Virtual Memory
On-device dispatch
Generic Address Space
Enhanced Image Support
C11 Atomics
Pipes
Android ICD

OpenCL C++ Kernel Language

Static subset of C++14

Templates and Lambdas

SPIR-V 1.2 with C++ support

SYCL 2.2 single source C++

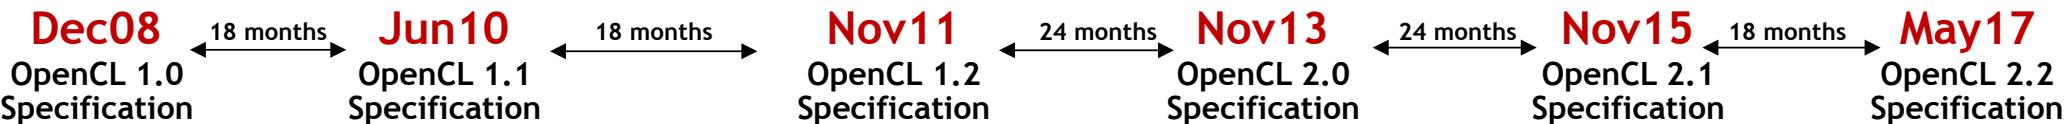
Pipes

Efficient device-scope
communication between kernels

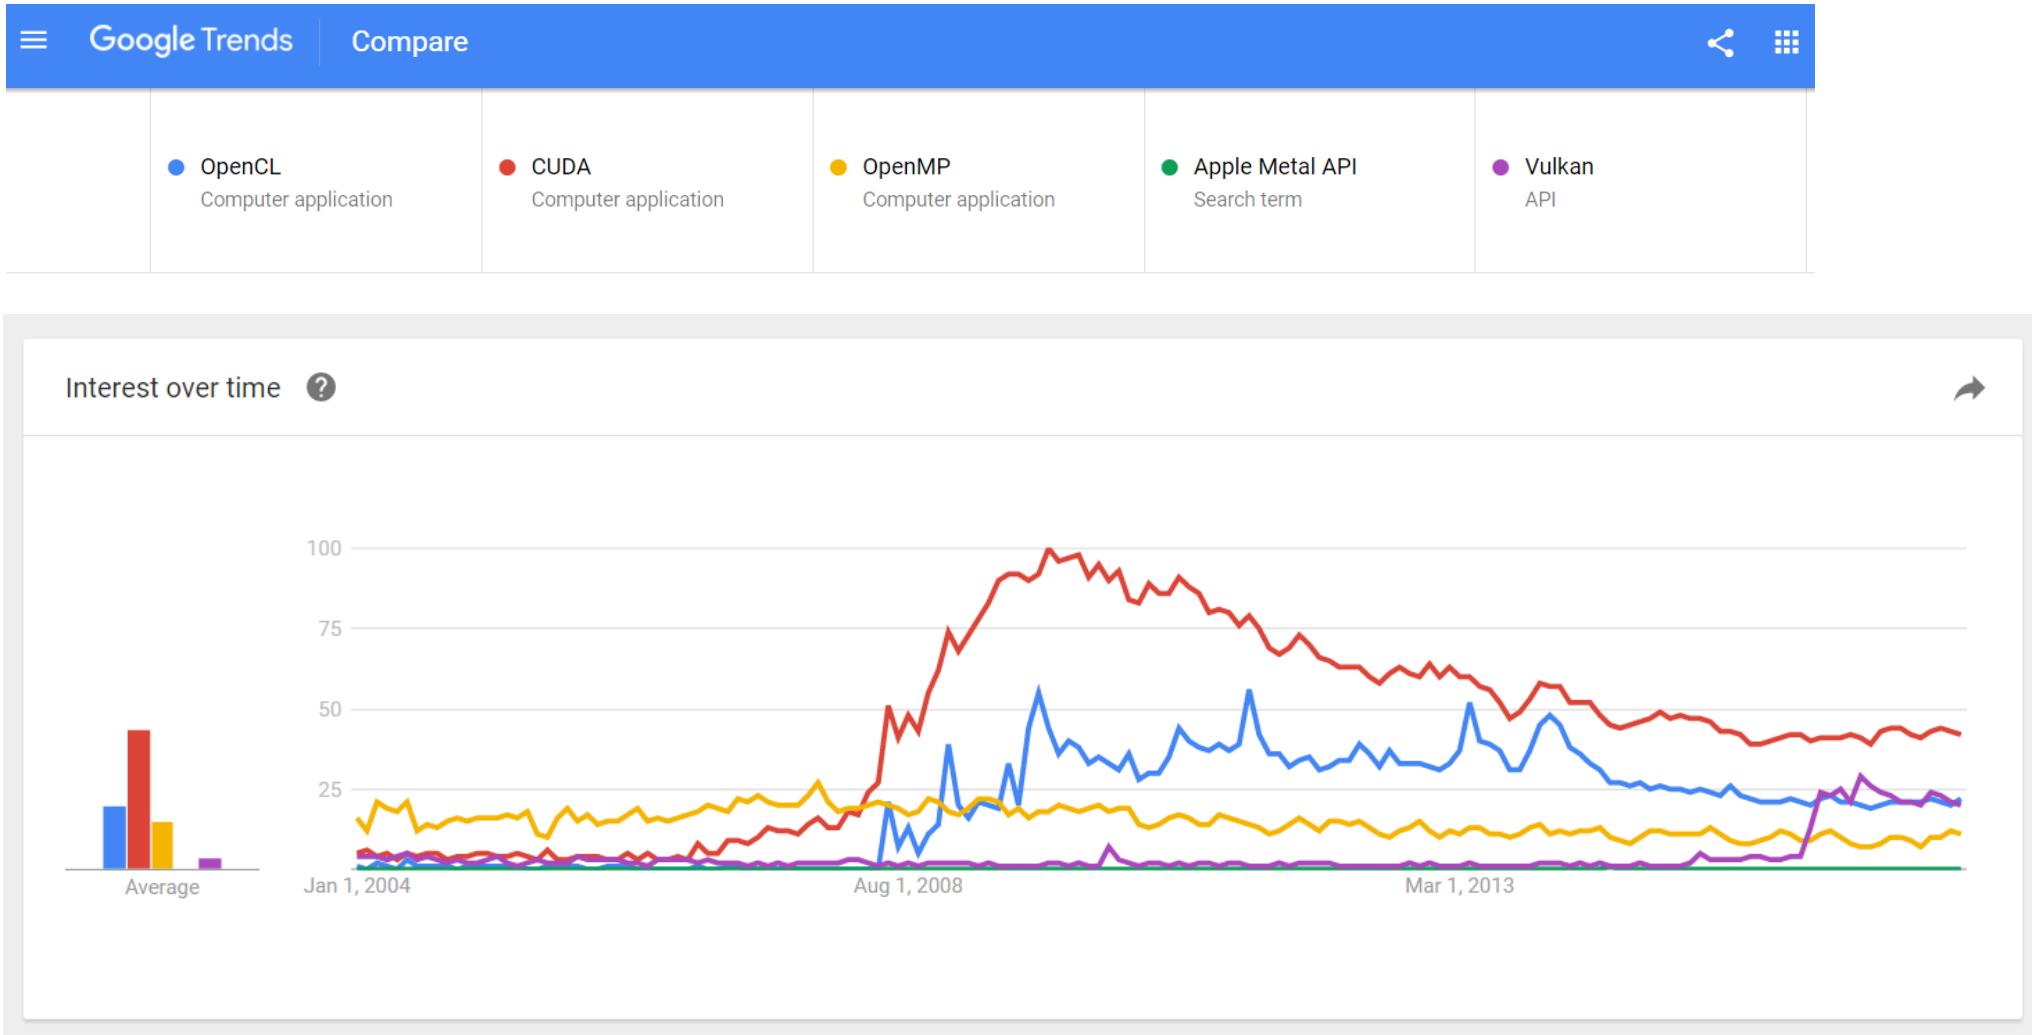
Multiple Code Generation
Optimizations

SPIR-V in Core
Subgroups into core
Subgroup query operations
clCloneKernel
Low-latency device
timer queries

LEARN



Google Trends

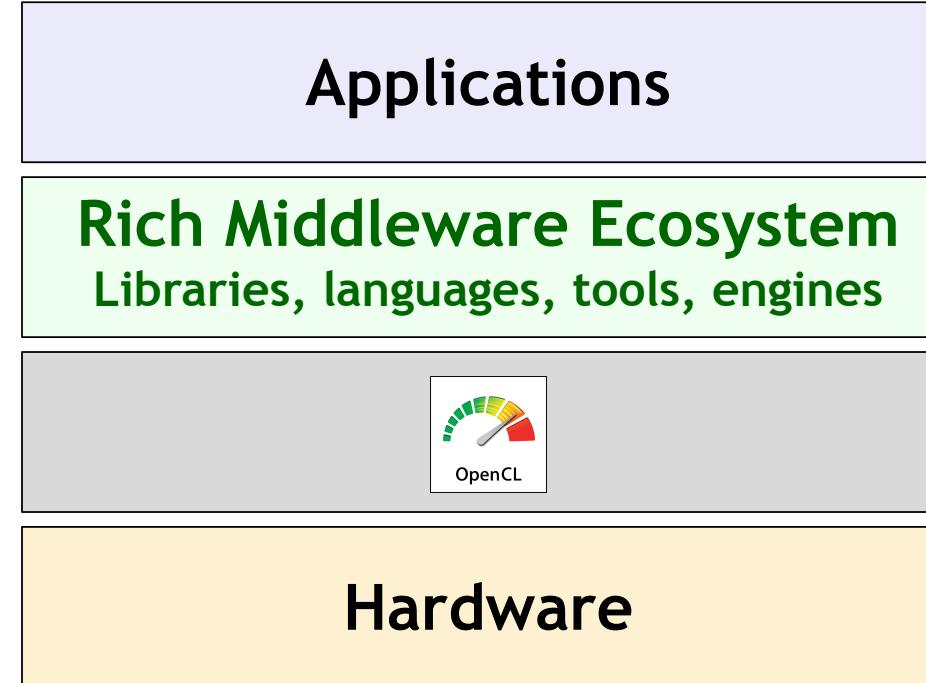


Embrace the Layered Ecosystem

OpenCL mixed
providing low-level
hardware access
with ‘ease-of-use’

Didn’t make it
clear that low-
level performance
portability is
impossible

Did not focus on
rapidly porting
efficient libraries



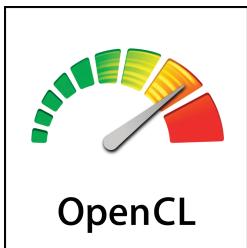
Middleware just
needs direct access
to hardware. Driver
should ‘get out of the
way’

Middleware can
provide ease of use

Middleware has the
system/domain
context to try to
provide performance
portability

- Run-time abstraction hardware is needed:
- Software vendors can't afford to port to every type/generation hardware
 - Hardware vendors want to keep innovating under an abstraction

Market Segments Need Deployment Flexibility



OpenCL has been over-monolithic

E.g. DSP inferencing should not be forced to ship IEEE FP32

Solution: feature sets - enabling toggling capabilities within a coherent framework without losing conformance

Desktop (actual and cloud)

Use cases: Video Image Processing, Gaming Compute, Rendering, Neural Network Training and Inferencing

Roadmap: Vulkan interop, dialable precision, pre-emption, collective programming and improved execution model, dynamic parallelism, pre-emption

Mobile

Use case: Photo and Vision Processing, Neural Network Inferencing

Roadmap: SVM, dialable precision for inference engine and pixel processing efficiency, pre-emption and QoS scheduling for power efficiency

HPC

Use case: Numerical Simulation, Neural Network Training, Virtualization

Roadmap: enhanced streaming processing, enhanced library support

FPGAs

Use cases: Network and Stream Processing

Roadmap: enhanced execution model, self-synchronized and self-scheduled graphs, fine-grained synchronization between kernels, DSL in C++

Embedded

Use cases: Vision, Signal and Pixel Processing, Neural Network and Inferencing

Roadmap: arbitrary precision for power efficiency, hard real-time scheduling, asynch DMA

Other Lessons

Lessons	How We Learned Them	How We Do Better!
Language flexibility is good! Enable language innovation!	OpenCL WG spent way too long designing OpenCL C and C++	Ingest SPIR-V! BUT Vendors need to support it!
Lack of tools and libraries	Assumption that the Working Group's job is done once the specification is shipped	'Hard launches' i.e. simultaneous availability of spec, libraries, implementations and engines
Needs to be adopted/available on key platforms	Apple are focused on Metal OpenCL/RenderScript Confusion NVIDIA not pushing to 2.0	Add value to key platforms and/or develop viable portability solutions
Middleware and application insights and prototyping are essential during standards design	The OpenCL Working Group has lacked active software developer participation	Encourage ISVs to join Khronos to help steer the industry! AND OpenCL Advisory Panels

Khronos Advisory Panels

The Working Group invites input and shares draft specifications and other WG materials



Members

Pay membership Fee
Sign NDA and IP Framework
Directly participate in working groups

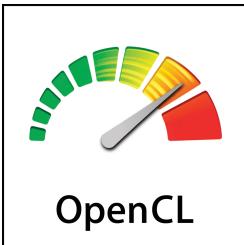
Advisors

Pay \$0
Sign Advisors Agreement = NDA and IP Framework
Provide requirements and feedback on specification drafts to the working group

Advisory Panel membership is
'By Invitation' and renewed annually.

No 'minimum workload' commitment - but we love input and feedback!
Please reach out if you wish to participate!

Requirements for ‘OpenCL Next’



Working Group Decision!
Converge with and leverage
Vulkan design!
Expand on Vulkan supported
processors types and
compute capabilities

Low-level explicit API as
Foundation of multi-layer ecosystem



Features set for
Market Deployment Flexibility



SPIR-V Ingestion for
Language flexibility



Widely Adopted
No market barriers to deployment



Installable tools architecture for
Development flexibility



Low-latency, multi-threaded dispatch
For fine-grained, high-performance



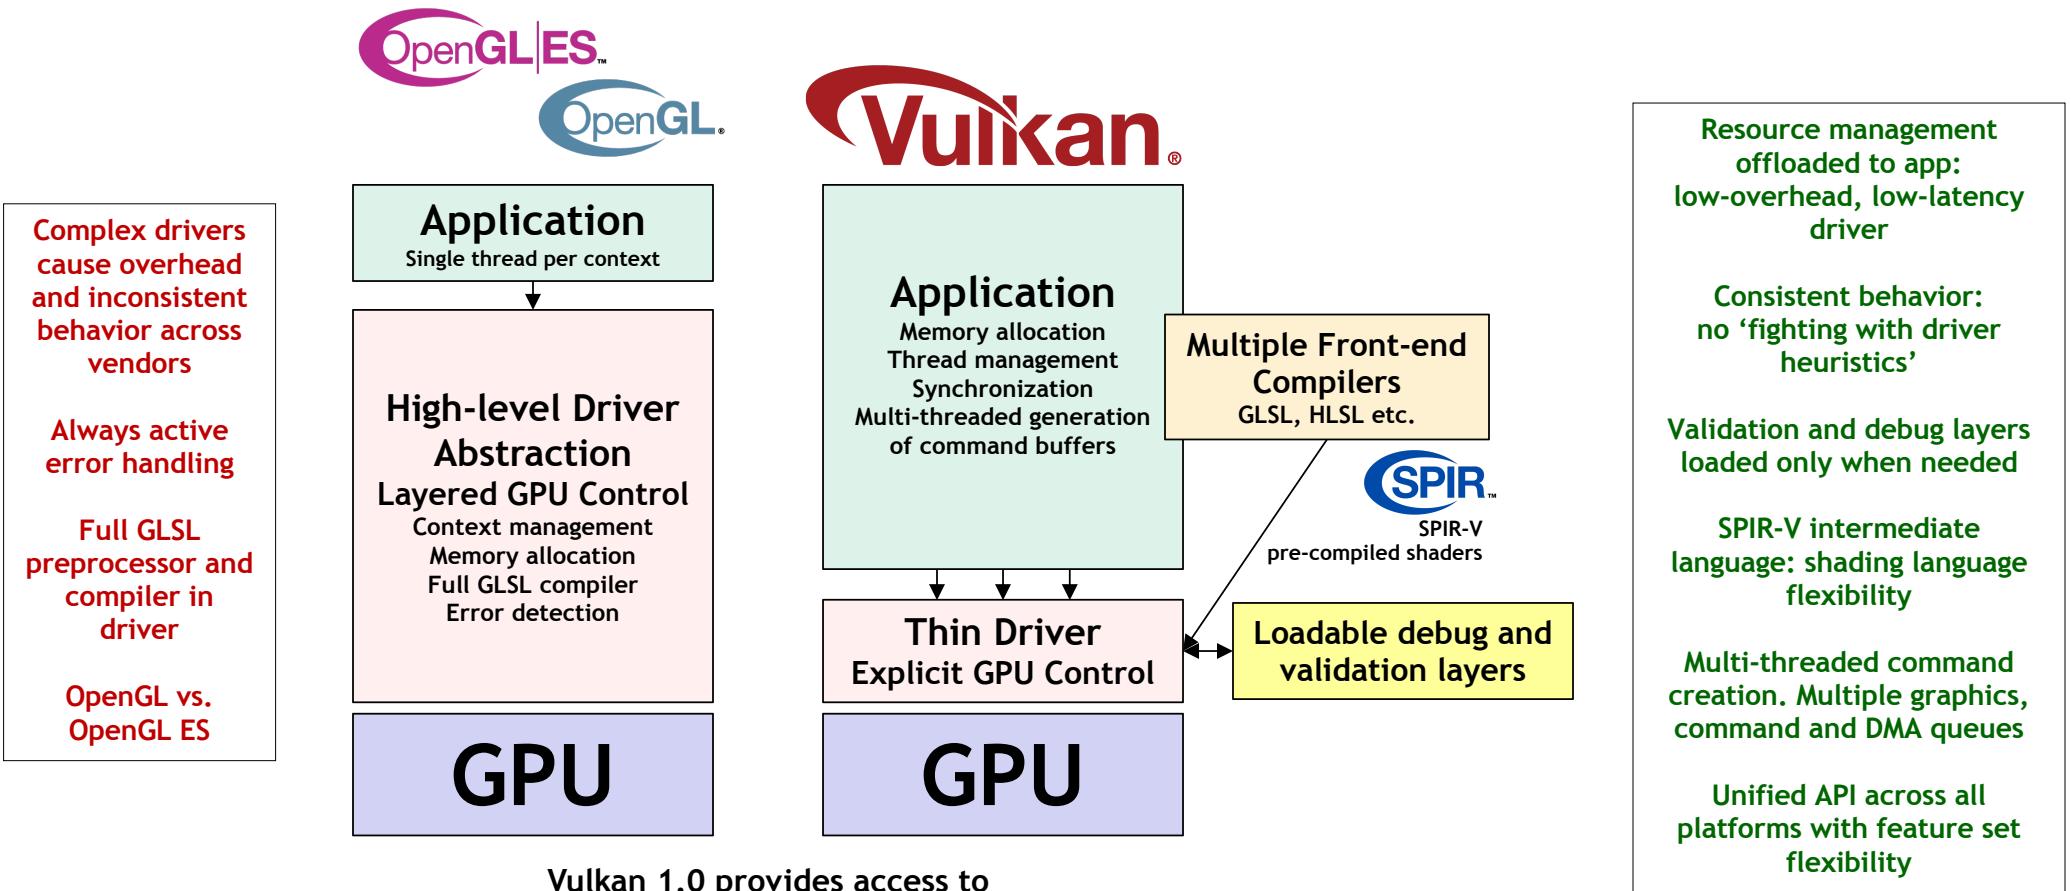
At least OpenCL 2.X-class
compute capabilities



Support for
diverse processor types



Vulkan Explicit GPU Control



Vulkan Adoption

All Major GPU Companies shipping Vulkan Drivers - for Desktop and Mobile Platforms



Mobile, Embedded and Console Platforms Supporting Vulkan



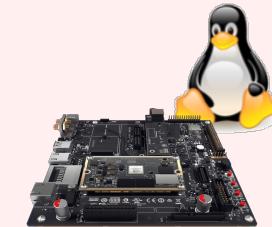
Android 7.0



Nintendo Switch



Android TV



Embedded Linux



Cross Platform



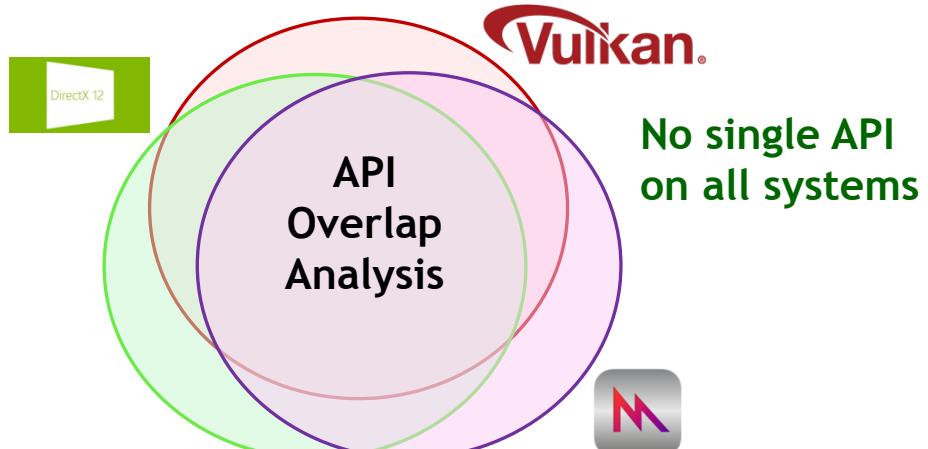
Windows 10



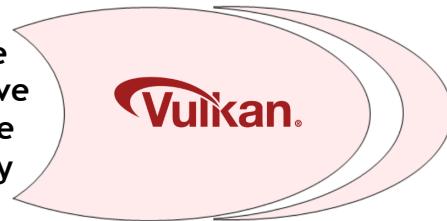
TIZEN™



GPU Portability - Call For Participation



Use Feature Sets to remove non-portable functionality

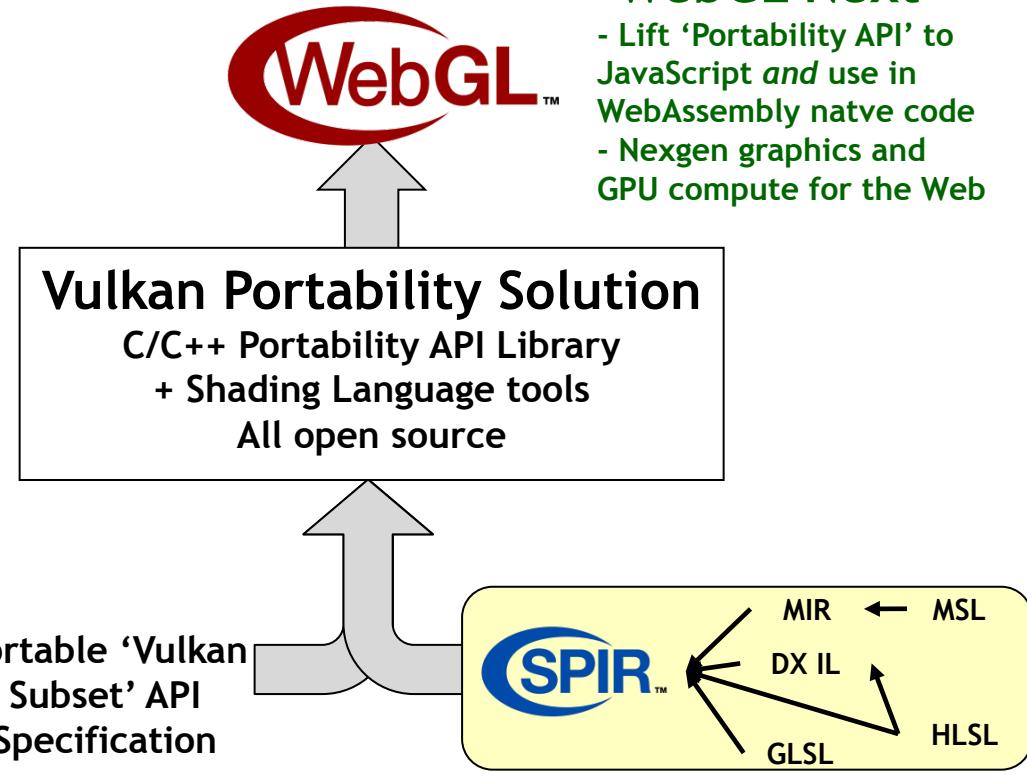


Use Extensions to add functionality e.g. security and robustness for the Web

No single API on all systems



Vulkan is non-proprietary and is already designed to be portable



Open source compilers/translators for shading and intermediate languages

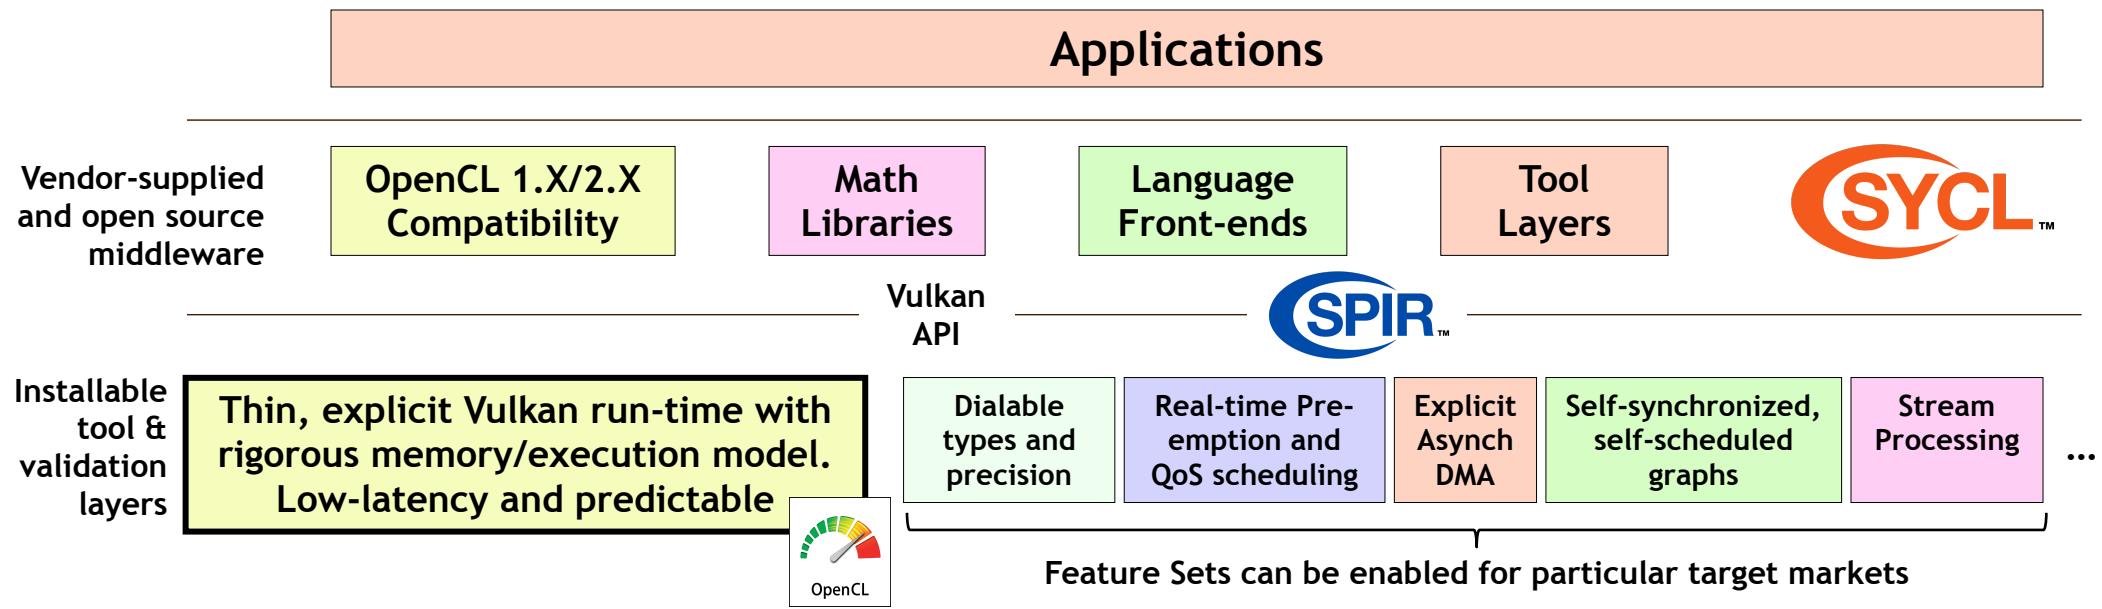
A Portability Solution needs to address APIs and shading languages

'WebGL Next'

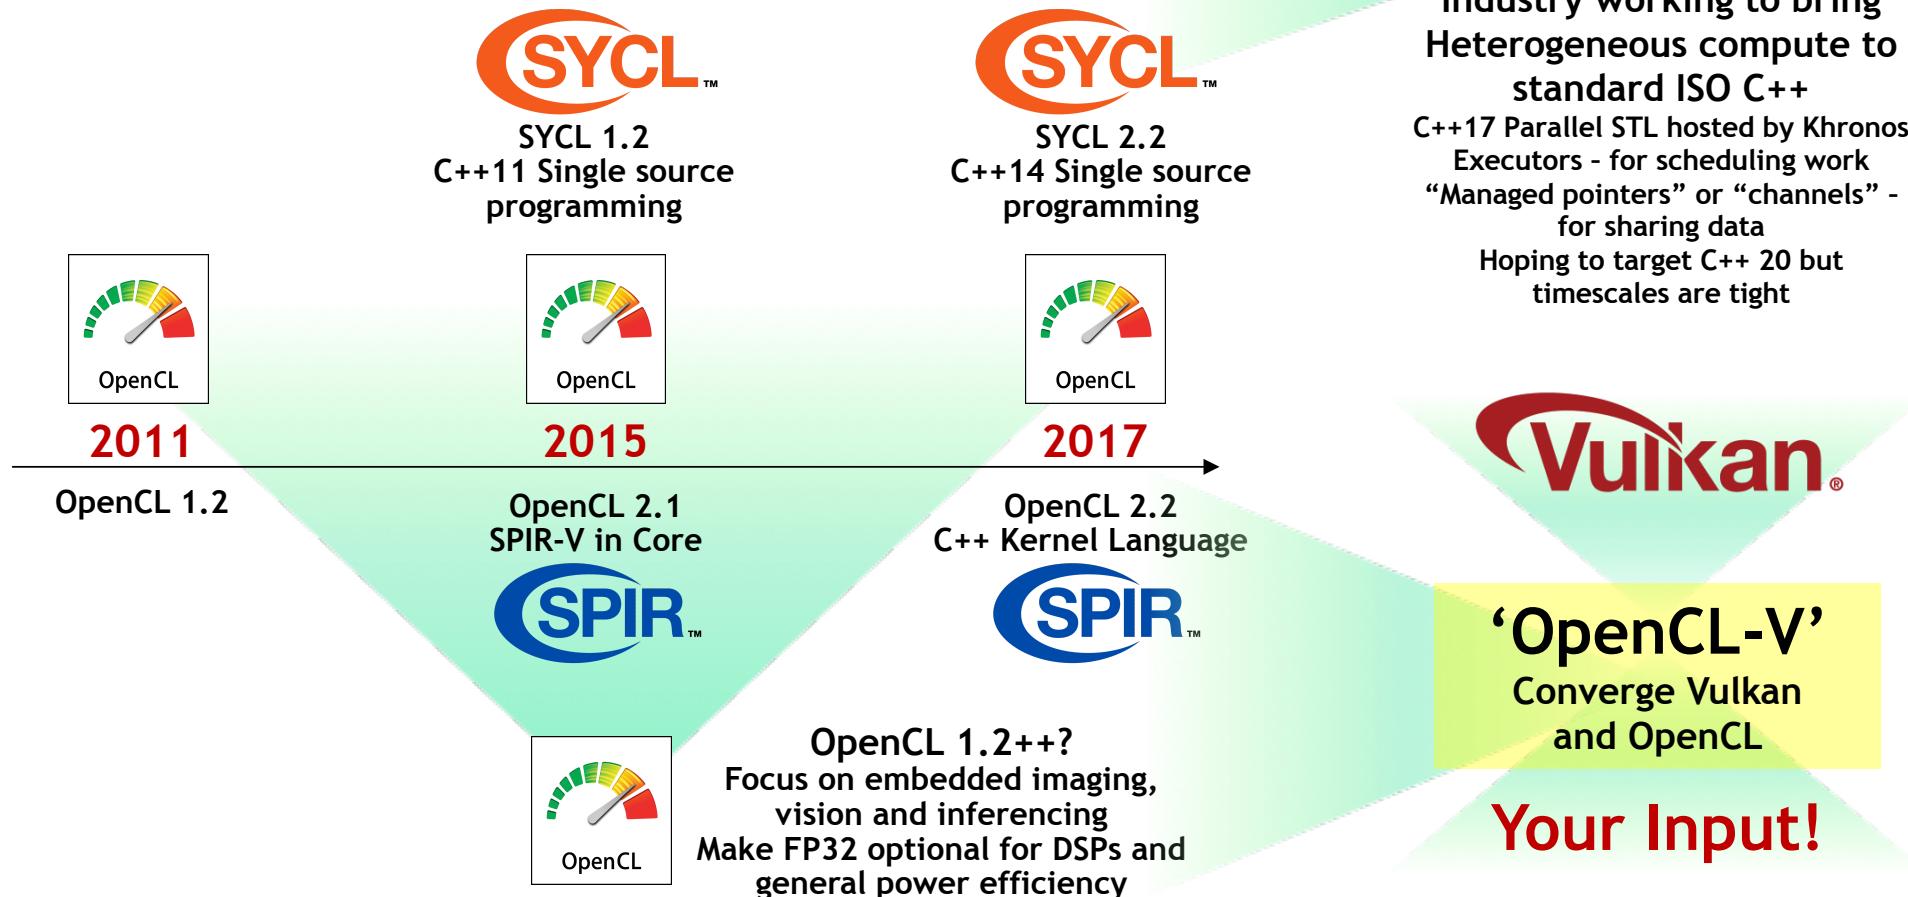
- Lift 'Portability API' to JavaScript *and* use in WebAssembly native code
- Nexgen graphics and GPU compute for the Web

‘OpenCL-V’ - OpenCL and Vulkan Convergence

- Converge OpenCL roadmap over time with Vulkan API and run-time
 - Support more processor types, e.g. DSPs and FPGAs (graphics optional)
- Layered ecosystem for backwards-compatibility and market flexibility
 - Feature sets for target market agility
- Single runtime stack for graphics *and* compute
 - Streamline development, adoption and deployment for the entire industry



OpenCL Evolution Discussions



Get Involved!

- OpenCL is driving to new level of community engagement
 - Learning from the Vulkan experience
 - We need to know what you need from OpenCL
 - IWOCL is the perfect opportunity to find out!
- Any company or organization is welcome to join Khronos
 - For a voice and a vote in any of these standards
 - www.khronos.org
- If joining is not possible - ask about the OpenCL Advisory Panel
 - Free of charge - enables design reviews, requirements and contributions
- Neil Trevett
 - ntrevett@nvidia.com
 - [@neilt3d](https://twitter.com/neilt3d)

